


# Chapter 9 – Input and Output Streams



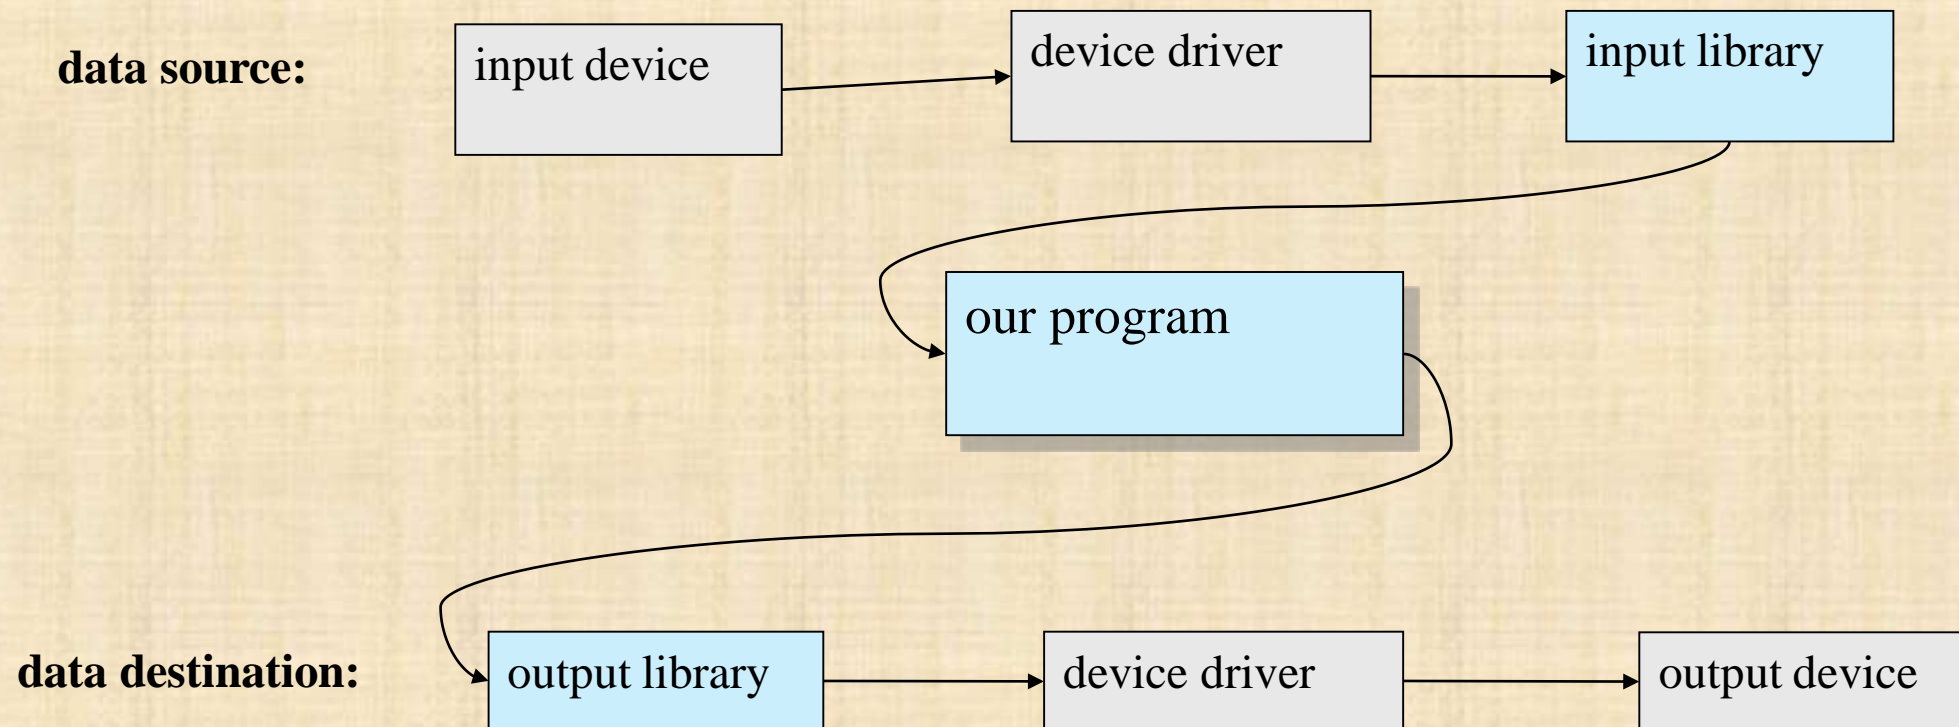
*Science is what we have learned  
about  
how to keep from fooling ourselves.  
– Richard P. Feynman*

# Overview

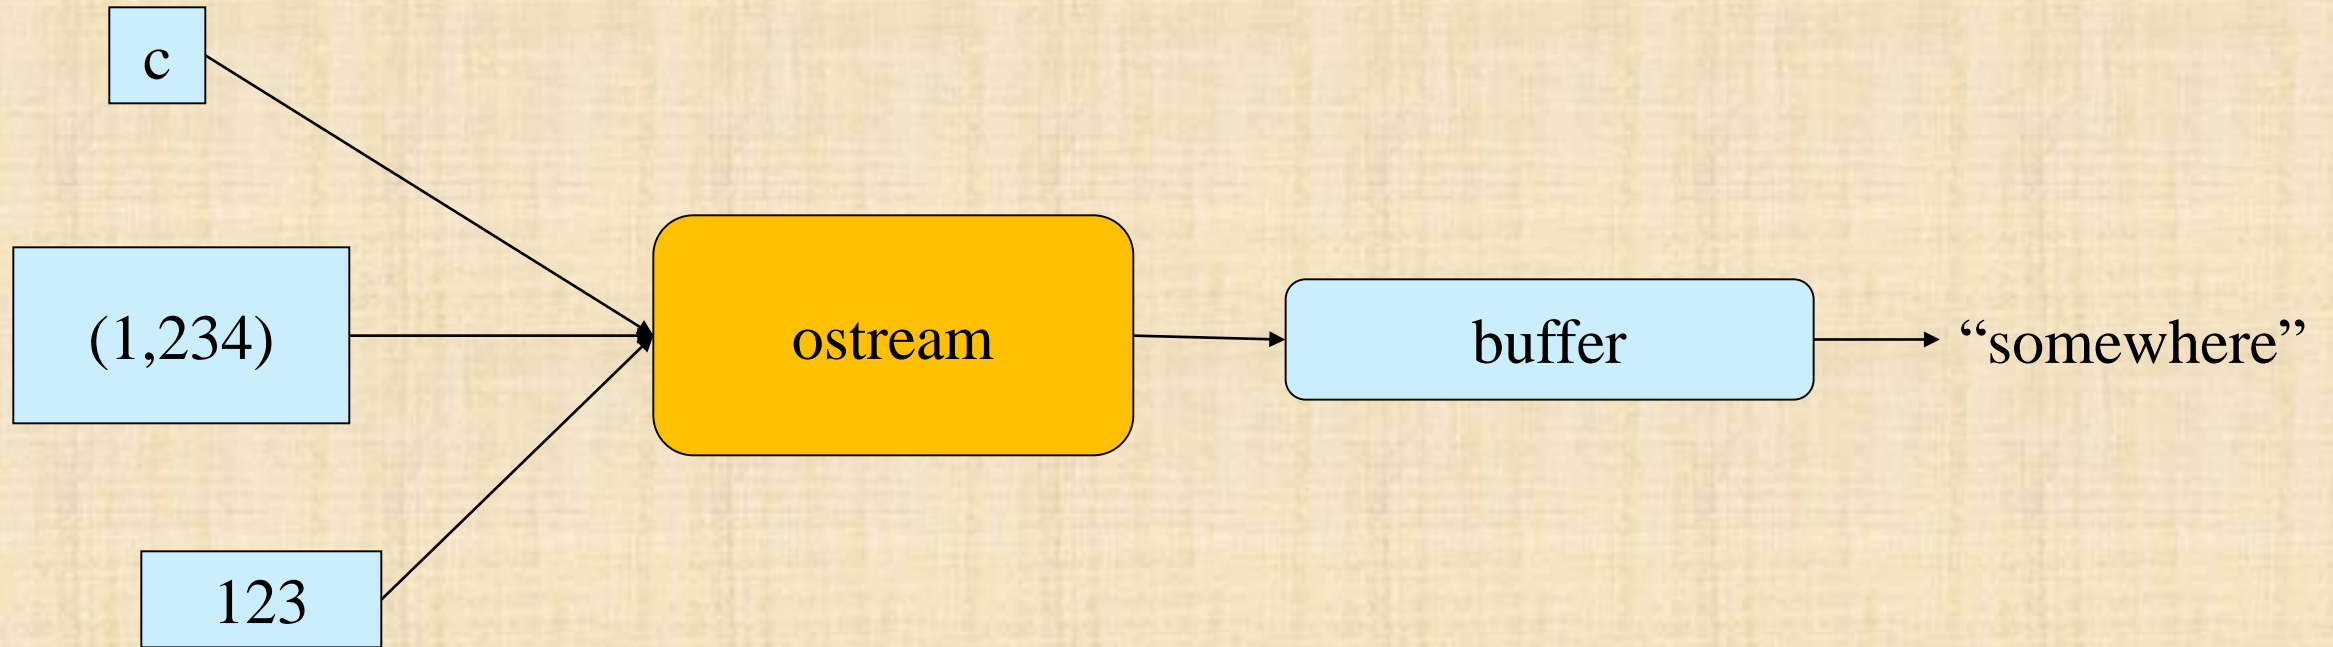
- Input and output
- The I/O stream model
- Files
  - Opening a file; Reading and writing a file
- I/O error handling
- Reading a single value
  - Breaking the problem into manageable parts; Separating dialog from function
- User-defined I/O
- Formatting
- Character I/O
- String streams



# Input and Output



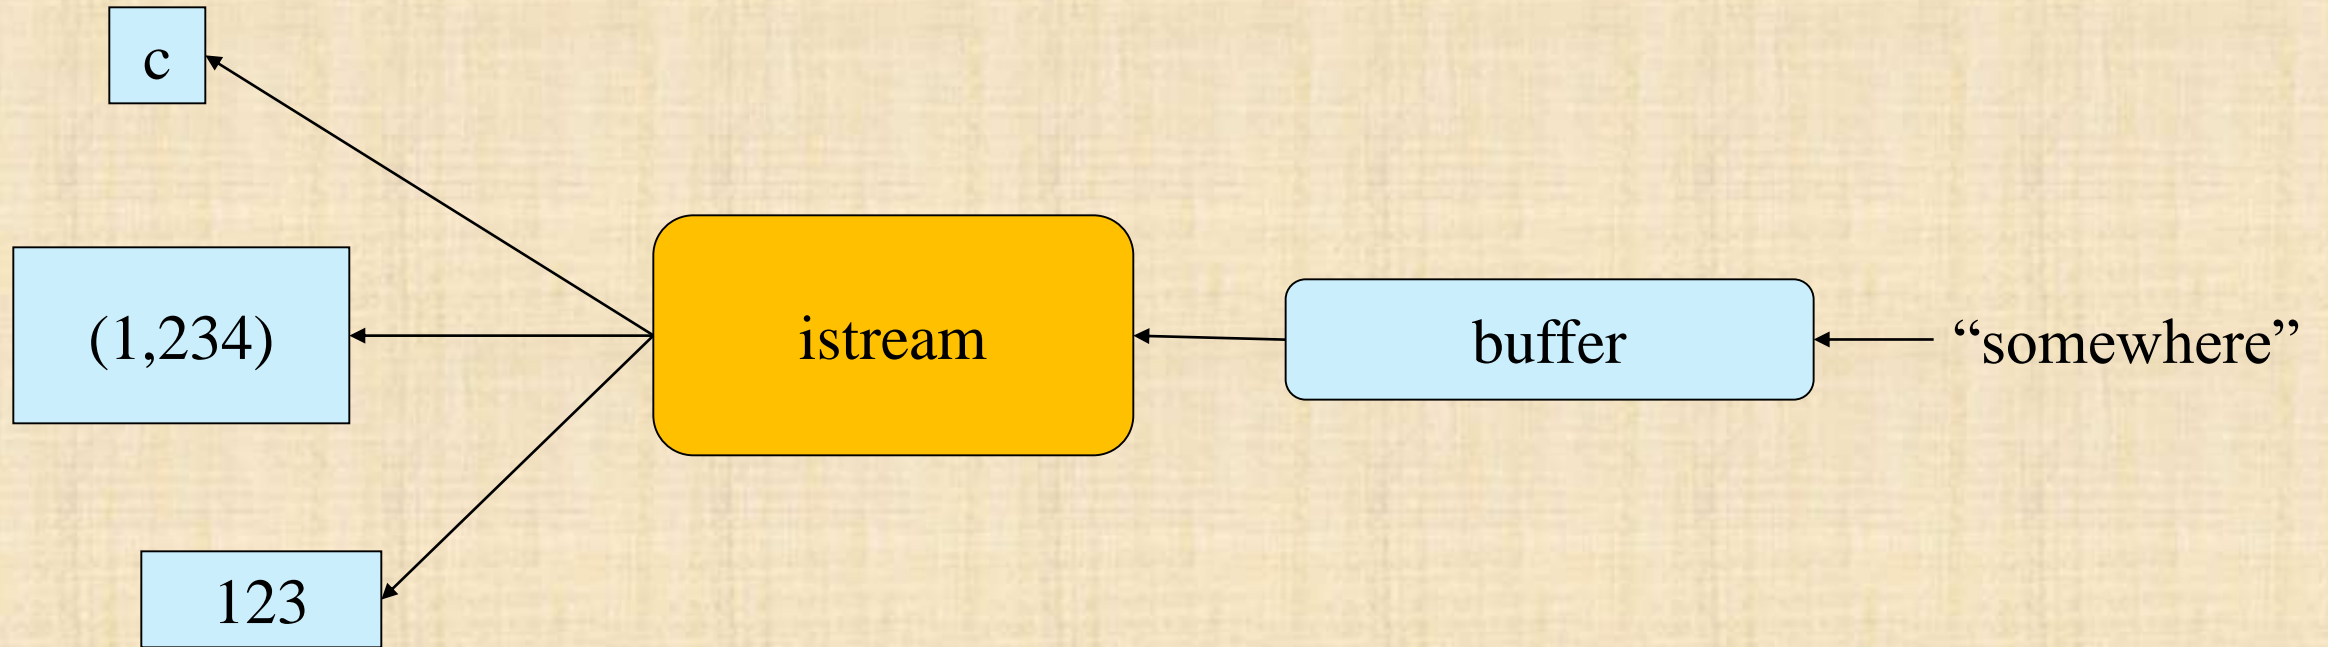
# The stream model



- An **ostream**
  - turns values of various types into character sequences
  - sends those characters somewhere
    - *E.g.*, console, file, main memory, another computer



# The stream model



- An **istream**
  - turns character sequences into values of various types
  - gets those characters from somewhere
    - *E.g.*, console, file, main memory, another computer

# The stream model

- Reading and writing
  - Of typed entities
    - `<<` (output) and `>>` (input) plus other operations
    - Type safe
    - Formatted
  - Typically stored (entered, printed, etc.) as text
    - But not necessarily
  - Extensible
    - You can define your own I/O operations for your own types
  - Stream oriented
    - A stream can be attached to any I/O or storage device
    - Some format properties can be set for all values sent to a stream



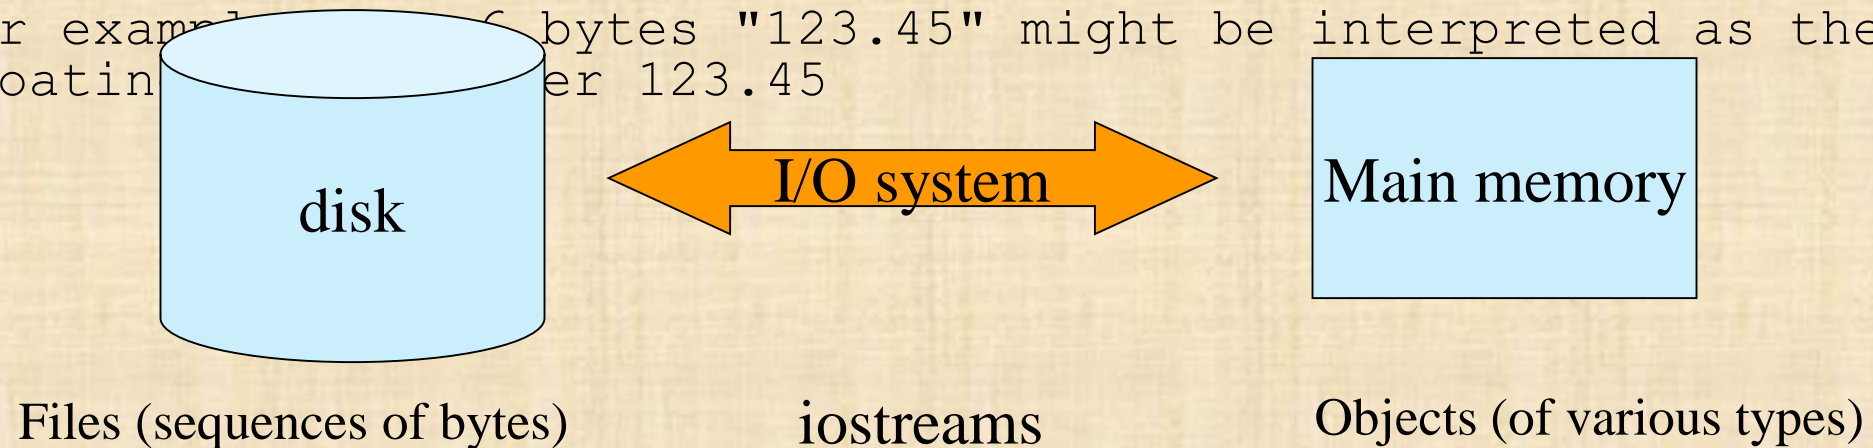
# Files

- We turn our computers on and off
  - The contents of our main memory is transient
- We like to keep our data
  - So we keep what we want to preserve on disks and similar permanent storage
- A file is a sequence of bytes stored in permanent storage
  - A file has a name
  - The data on a file has a format
- We can read/write a file if we know its name and format

# A file



- At the fundamental level, a file is a sequence of bytes numbered from 0 upwards
- Other notions can be supplied by programs that interpret a "file format"
  - For example, 6 bytes "123.45" might be interpreted as the floating-point number 123.45





# Files

- To read a file
  - We must know its name
  - We must open it (for reading)
    - An attempt to open a file may fail
  - Then we can read
  - Then we must close it
    - That is typically done implicitly
- To write a file
  - We must name it
  - We must open it (for writing)
    - Or create a new file of that name
    - An attempt to open a file may fail
  - Then we can write it
  - We must close it
    - That is typically done implicitly

# Opening a file for reading

- To read a file
  - We must know its name
  - We must open it (for reading)
    - An attempt to open a file may fail
  - Then we can read
  - Then we must close it
    - That is typically done implicitly

```
cout << "Please enter input file name: ";  
string iname;  
cin >> iname;  
ifstream ist {iname}; // ifs is an input stream from  
named iname  
if (!ist)  
    error("can't open input file ", iname);
```



# Opening a file for writing

- To write a file
  - We must name it
  - We must open it (for writing)
    - Or create a new file of that name
    - An attempt to open a file may fail
  - Then we can write it
  - We must close it
    - That is typically done implicitly

```
cout << "Please enter name of output file: ";  
string oname;  
cin >> oname;  
ofstream ofs {oname};    // ofs is an output stream from  
                           a file name oname  
if (!ofs)  
    error("can't open output file ", oname);
```

# Streams have destructors

- Use constructors and destructors to control lifetime

```
int main()
{
    string iname;
    // ...
    istream ist {iname};
    // ...
    string oname;
    // ...
    ostream {oname};
    // ...
} // here, files are implicitly closed, I/O buffers and strings
are given back to the free store
```



# Reading from a file

- Suppose a file contains a sequence of pairs representing hours and temperature readings

0 60.7

1 60.6

2 60.3

3 59.22

- The hours are numbered 0..23
- No further format is assumed
  - Maybe we can do better than that (but not just now)
- Termination
  - Reaching the end of file terminates the read
  - Anything unexpected in the file terminates the read
    - *E.g., q*

# Reading a file

```
struct Reading { // a temperature reading
    int hour;    // hour after midnight [0:23]
    double temperature;
};

vector<Reading> temps; // create a vector to store the readings

int hour;
double temperature;
while (ist >> hour >> temperature) {                                // read
    if (hour < 0 || 23 < hour)                                         // check (always
check validity of input)
        error("hour out of range");
    temps.push_back( Reading{hour, temperature} );                  // store
}
```



# I/O error handling

- Sources of errors
  - Human mistakes
  - Files that fail to meet specifications
  - Specifications that fail to match reality
  - Programmer errors
  - Etc.
- **iostream** reduces all errors to one of four states
  - **good()** *// the operation succeeded*
  - **eof()** *// we hit the end of input ("end of file")*
  - **fail()** *// something unexpected happened*
  - **bad()** *// something unexpected and serious happened*

# Sample integer read "failure"

- Ended by "terminator character" - program to deal with that
  - 1 2 3 4 5 \*
  - State is **fail()**
- Ended by format error - program to deal with that
  - 1 2 3 4 5.6
  - State is **fail()**
- Ended by "end of file" - often the desired end
  - 1 2 3 4 5 end of file
  - 1 2 3 4 5 Control-Z (Windows)
  - 1 2 3 4 5 Control-D (Unix)
  - State is **eof()**
- Something really bad - you usually can't recover from that
  - Disk format error
  - State is **bad()**



# I/O error handling

```
void fill_vector(istream& ist, vector<int>& v, char terminator)
    // read integers from ist into v until we reach eof() or
    terminator
{
    for (int i; ist >> i; )    // read until "some failure"
        v.push_back(i);    // store in v
    if (ist.eof()) return;    // fine: we found the end of
    file
    if (ist.bad()) error("ist is bad");    // stream corrupted; let's
    get out of here!

    if (ist.fail()) {    // clean up the mess as best we can and report
    the problem
        ist.clear();    // clear stream state, so that we can look
    for terminator
        char c;
        ist >> c;    // read a character, hopefully terminator
        if (c != terminator) {    // unexpected character
            ist.unget();    // put that character back
            ist.clear(ios base::failbit);    // set the state
```

# Throw an exception for `bad()`

- Often, we cannot recover from **bad**
  - E.g., disk read error, unhandled network failure

*// How to make **ist** throw if it goes **bad**:*

```
ist.exceptions(ist.exceptions() | ios_base::badbit) ;
```

*// can be read as*

*// "set **ist**'s exception mask to whatever it was plus  
badbit"*

*// or as "throw an exception if the stream goes bad"*

Given that, we can simplify our input loop by no longer checking for **bad**



## Simplified input loop

```
void fill_vector(istream& ist, vector<int>& v, char
    terminator)
{
    // read integers from ist into v until we reach
    eof() or terminator
    for (int i; ist >> i; )
        v.push_back(i);
    if (ist.eof()) return; // fine: we found the end of file

    // not good() and not bad() and not eof(), ist must be
    fail()
    ist.clear(); // clear stream state
    char c;
    ist >> c; // read a character, hopefully terminator
    if (c != terminator) { // ouch: not the terminator, so we
        must fail
        ist.unget(); // maybe my caller can use that
        character
        ist.clear(ios_base::failbit); // set the state back to
        fail()
    }
```

## Throw an exception for `bad()`

- Often, we don't want to try to recover from **fail**
  - E.g., file-read format error

```
ist.exceptions(ist.exceptions() | ios_base::badbit  
() | ios_base::failbit);
```

Given that, we can simplify our input loop to the simple and obvious:

```
void fill_vector(istream& ist, vector<int>& v)  
    // read integers from ist into v until we reach eof()  
{  
    for (int x; ist >> x; )  
        v.push_back(x);  
}
```



# Reading a single value

```
// first simple and flawed attempt:
cout << "Please enter an integer in the range 1 to 10
      (inclusive):\n";
int n = 0;
while (cin>>n) {                // read
    if (1<=n && n<=10) break;    // check range
    cout << "Sorry, "
         << n
         << " is not in the [1:10] range; please try
again\n";
}

// use n here
```

- Three kinds of problems are possible
  - the user types an out-of-range value
  - getting no value (end of file)
  - the user types something of the wrong type (here, not

# Reading a single value

- What do we want to do in those three cases?
  - handle the problem in the code doing the read?
  - throw an exception to let someone else handle the problem (potentially terminating the program)?
  - ignore the problem?
- Reading a single value
  - Is something we often do many times
  - We want a solution that's very simple to use



# Handle everything: What a mess!

```
cout << "Please enter an integer in the range 1 to 10
(inclusive):\n";
int n = 0;
while (cin >> n) {
    if (cin) { // we got an integer; now check it:
        if (1<=n && n<=10) break;
        cout << "Sorry, " << n << " is not in the [1:10] range; please
try again\n";
    }
    else if (cin.fail()) { // we found something that wasn't an
integer
        cin.clear(); // we'd like to look at the characters
        cout << "Sorry, that was not a number; please try again\n";
        for (char ch; cin>>ch && !isdigit(ch); ) /* nothing */ ; //
throw away non-digits
        if (!cin) error("no input"); //
we didn't find a digit: give up
        cin.unget(); // put the digit back, so that we can read the
number
    }
else
```

# The mess: trying to do everything at once

- Problem: We have all mixed together
  - reading values
  - prompting the user for input
  - writing error messages
  - skipping past "bad" input characters
  - testing the input against a range
- Solution: Split it up into logically separate parts



# What do we want?

- What logical parts do we want?
  - `int get_int(int low, int high);`    *// read an int in [**low**..**high**] from **cin***
  - `int get_int();`                    *// read an int from **cin** so that we can check the range **int***
  - `void skip_to_int();`                *// we found some "garbage" character so skip until we find an int*
- Separate functions that do the logically separate actions

# Skip "garbage"

```
void skip_to_int()
{
    if (cin.fail()) {                // we found something that wasn't an
        integer                      integer
        cin.clear();                // we'd like to look at the characters
        for(char ch; cin>>ch; ) {    // throw away non-digits
            if (isdigit(ch) || ch=='-') {
                cin.unget();          // put the digit back, so that we
                can read the number
                return;
            }
        }
    }
    error("no input");              // eof or bad: give up
}
```



## Get (any) integer

```
int get_int()
{
    int n = 0;
    while (true) {
        if (cin >> n)
            return n;
        cout << "Sorry, that was not a number; please try again\n";
        skip_to_int();
    }
}
```

# Get integer in range

```
int get_int(int low, int high)
{
    cout << "Please enter an integer in the range "
          << low << " to " << high << " (inclusive):\n";
    while (true) {
        int n = get_int();
        if (low<=n && n<=high)
            return n;
        cout << "Sorry, "
              << n << " is not in the [" << low << ':' <<
high
              << "]" range; please try again\n";
    }
}
```



# Use

```
int n = get_int(1,10);  
cout << "n: " << n << endl;
```

```
int m = get_int(2,300);  
cout << "m: " << m << endl;
```

- Problem:
  - The “dialog” is (still) built into the read operations

What do we ***really*** want?

```
// parameterize by integer range and "dialog"
```

```
int strength = get_int(1, 10,  
                      "enter strength",  
                      "Not in range, try again");  
cout << "strength: " << strength << endl;  
  
int altitude = get_int(0, 50000,  
                      "please enter altitude in feet",  
                      "Not in range, please try again");  
cout << "altitude: " << altitude << "ft. above sea level\n";
```

- That's often the really important question
- Ask it repeatedly during software development
- As you learn more about a problem and its solution, your answers improve



# Parameterize

```
int get_int(int low, int high, const string& greeting, const
string& sorry)
{
    cout << greeting << ": [" << low << ':' << high << "]\n";
    while (true) {
        int n = get_int();
        if (low<=n && n<=high) return n;
        cout << sorry << ": [" << low << ':' << high << "]\n";
    }
}
```

- Incomplete parameterization: **get\_int()** still “blabbers”
  - “utility functions” should not produce their own error messages
  - Serious library functions do not produce error messages at all
    - They throw exceptions (possibly containing an error message)

# User-defined output: operator<<()

- Usually trivial

```
ostream& operator<<(ostream& os, const Date& d)
{
    return os << '(' << d.year()
        << ',' << d.month()
        << ',' << d.day() << ')';
}
```

- We often use several different ways of outputting a value
  - Tastes for output layout and detail vary



# Use

```
void do_some_printing(Date d1, Date d2)
{
    cout << d1;                // means operator<<(cout,d1) ;

    cout << d1 << d2;          // means (cout << d1) << d2;
                                // means (operator<<(cout,d1)) << d2;
                                // means operator<<((operator<<(cout,d1)), d2)
    ;
}
```

## User-defined input: operator>>()

- *//* Input is usually messier than input: formatting and errors

```
istream& operator>>(istream& is, Date& dd)
    // Read date in format: ( year , month , day )
{
    int y, d, m;
    char ch1, ch2, ch3, ch4;
    is >> ch1 >> y >> ch2 >> m >> ch3 >> d >> ch4;
    if (!is) // we didn't get our
        values, so just leave
        return is;
    if (ch1!='(' || ch2!=',' || ch3!=',' || ch4!=')') { // oops:
        format error
        is.clear(ios_base::failbit); // something wrong: set state to
        fail()
        return is; // and leave
    }
    dd = Date{y, Month(m), d}; // update dd
```



# Character I/O

- Sometimes, we have to drop down a level and read and look at individual characters

- E.g., to tokenize **1+4\*x<=y/z\*5**

```
for (char ch; cin.get(ch); ) {           // get(ch) reads the next character
    into ch; doesn't skip whitespace

    if (isspace(ch)) {
        // do nothing; i.e., skip whitespace (e.g. space or tab)
    }
    else if (isdigit(ch)) {
        // .. read a number ...
    }
    else if (isalpha(ch)) {
        // ... read an identifier ...
    }
    else {
        // ... deal with operators ...
    }
}
```

# Character classification

- Useful (and fast) standard-library functions for dealing with characters

Character classification	
<code>isspace(c)</code>	Is <code>c</code> whitespace (' ', '\t', '\n', etc.)?
<code>isalpha(c)</code>	Is <code>c</code> a letter ('a'..'z', 'A'..'Z') (note: not '_')?
<code>isdigit(c)</code>	Is <code>c</code> a decimal digit ('0'..'9')?
<code>isxdigit(c)</code>	Is <code>c</code> a hexadecimal digit (decimal digit or 'a'..'f' or 'A'..'F')?
<code>isupper(c)</code>	Is <code>c</code> an uppercase letter?
<code>islower(c)</code>	Is <code>c</code> a lowercase letter?
<code>isalnum(c)</code>	Is <code>c</code> a letter or a decimal digit?
<code>isctrl(c)</code>	Is <code>c</code> a control character (ASCII 0..31 and 127)?
<code>ispunct(c)</code>	Is <code>c</code> not a letter, digit, whitespace, or invisible control character?
<code>isprint(c)</code>	Is <code>c</code> printable (ASCII ' '.. '~')?
<code>isgraph(c)</code>	Is <code>isalpha(c)</code> or <code>isdigit(c)</code> or <code>ispunct(c)</code> (note: not space)?

Character case	
<code>x=toupper(c)</code>	<code>x</code> becomes <code>c</code> or <code>c</code> 's uppercase equivalent
<code>x=tolower(c)</code>	<code>x</code> becomes <code>c</code> or <code>c</code> 's lowercase equivalent



# String streams – using strings (in memory) as I/O sources and targets

```
Point get_coordinates(const string& s)           // extract {x,y} from "(x,y)"
{
    istringstream is {s}; // make a stream so that we can read from s
    Point xy;
    char left_paren, ch, right_paren;
    is >> left_paren >> xy.x >> ch >> xy.y >> right_paren;
    if (!is || left_paren != '(' || ch != ',' || right_paren != ')')
        error("format error: ",s);
    return xy;
}

auto c1 = get_coordinates("(2,3)");
auto c2 = get_coordinates("( 200, 300) ");
auto c3 = get_coordinates("100,400");           // will call error()
```

# Formatting

- An integer can be printed in different bases
  - Decimal, hexadecimal (base 16), octal (base 8), and binary (base 2)
- A floating-point value can be printed in one of 4 formats
  - **fixed** (e.g., **12.67**), scientific (**e.g., 1.234e7**),
  - **defaultfloat** (what fits best in a field), **hexfloat** (hexadecimal mantissa and exponent)
- You can choose the **precision** used for writing a number
- You can choose the **width** of a field for a number (the number of character positions used)
- You can choose the placement of a number within a field (**adjustfield**) (**left, right, internal**)
- See PPP3§9.10 for some details, and cppreference for many details



# Formatting

- **printf()** is arguably that most popular C function, and a reason for C's success
- But it is not type safe

```
printf("an int %g and a string '%s'\n", "Hello!", 123);    // oops!
```
- Nor is it easily extensible

```
printf("a Point %P\n", Point{100,200});    // oops!
```
- C++ now have a printf()-like format() without the opportunities for type errors

```
cout << format("an int {} and a string '{}'\n", "Hello!", 123);    // print the arguments according to their type
```

  - And a (less obvious) mechanism for extensibility

```
cout << format("a Point %P\n", Point{100,200});    // works if you have defined %P to refer to Point
```
- See PPP3§9.10.6 for some details, and cppreference for many details

# Next Lecture

Graphical output