


Chapter 15 – **vector** and Free Store



*Use **vector** as the default!*
– Alex Stepanov

Overview: evolving a **vector** type

(and spanning the low-level/high-level language gap in the process)

- Chapter 15
 - Dealing with “raw” memory: pointers and free store
 - Destructors
- Chapter 16
 - Arrays and pointers
 - Pointers and references
 - C-style strings
 - Alternatives to low level facilities: **span**, **array**, and **string**
- Chapter 17
 - Copying and moving
 - Essential operations for managing object lifecycles
- Chapter 18
 - Templates and generic programming
 - Exceptions and scope-based resource management and error handling (RAII)
 - Resource management pointers: **unique_ptr** and **shared_ptr**

Vector

- **vector** is the most useful container
 - ISO standard
 - Simple to use
 - Compactly stores elements of a given type
 - Efficient access
 - Expands to hold any number of elements
 - Optionally range-checked access
 - the PPP version is range checked
- How is that done?
 - That is, how is **vector** implemented?
 - We'll answer that gradually, feature after feature
- Prefer **vector** for storing elements unless there's a good reason not to

Building from the ground up

- The hardware provides memory and addresses
 - Low level
 - Untyped
 - Fixed-sized chunks of memory
 - No checking
 - As fast as the hardware architects can make it
- The application builder needs something like a **vector**
 - Higher-level operations
 - Type checked
 - Size varies (as we get more data)
 - Run-time range checking
 - Close to optimally fast

Building from the ground up

- At the lowest level, close to the hardware, life's simple and brutal
 - You have to program everything yourself
 - You have no type checking to help you
 - Run-time errors are found when data is corrupted or the program crashes
- We want to get to a higher level as quickly as we can
 - To become productive and reliable
 - To use a language "fit for humans"
- Chapters 15-18 basically show all the steps needed
 - The alternative to understanding is to believe in "magic"
 - The techniques for building **vector** are the ones underlying all higher-level work with data structures

Vector

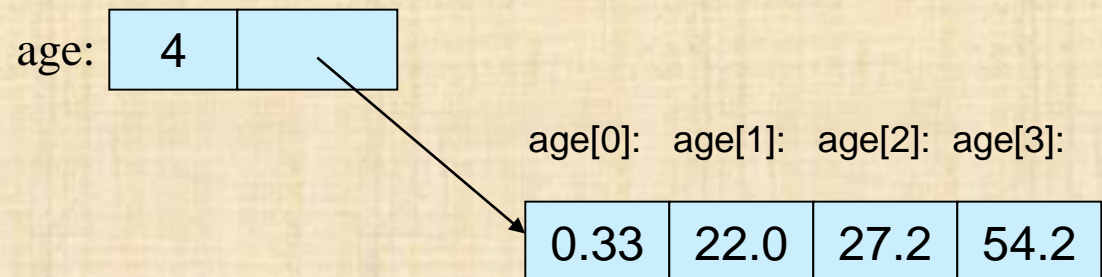
- A **vector**

- Our **Vector** will gradually be improved to approximate the standard **vector**
- Can hold an arbitrary number of elements
 - Up to whatever physical memory and the operating system can handle
- That number can vary over time
 - E.g. by using `push_back()`

- Example

```
Vector<double> age(4);
```

```
age[0]=.33; age[1]=22.0; age[2]=27.2; age[3]=54.2;
```



Vector

```
class Vector {    // a very simplified vector of doubles (like
    vector<double>)
    int sz;           // the number of elements ("the size")
    double* elem;     // pointer to the first element
public:
    Vector(int s);    // constructor: allocate s elements,
    let elem point to them
    int size() const { return sz; }    // the current size
};
```

- ***** means "pointer to" so **double*** is a "pointer to **double**"

- What is a "pointer"?
- How do we make a pointer "point to" elements?
- How do we "allocate" elements?


Pointer values

- Pointer values are memory addresses
 - Think of them as a kind of integer values
 - The first byte of memory is 0, the next 1, and so on
 - A pointer **p** can hold the address of a memory location



- A pointer points to an object of a given type
 - E.g., a **double*** points to a **double**, not to a **string**
- A pointer's type determines how the memory referred to by the pointer's value is used
 - E.g., what a **double*** points to can be added but not, say, concatenated

The computer's memory

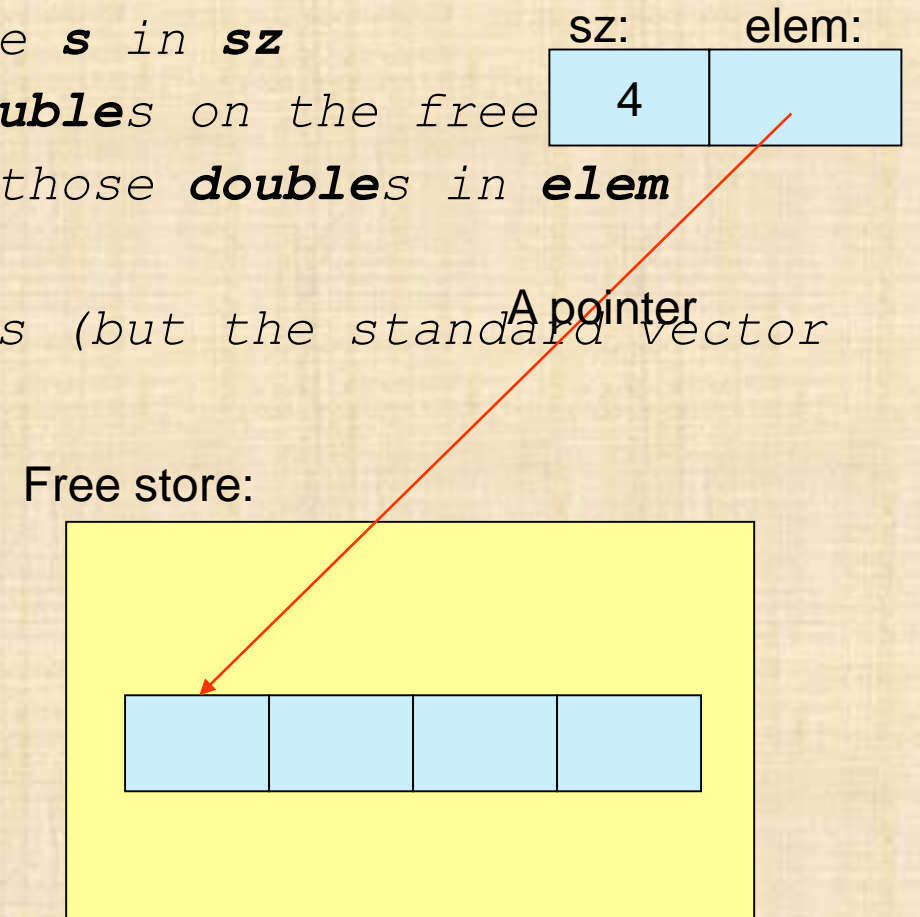


- As a program sees it
 - Local variables “live on the stack” (including function arguments)
 - Global variables are “static data”
 - The executable code is in “the code section”

Vector (constructor)

```
Vector::Vector(int s)    // Vector's constructor
    :sz(s),              // store the size s in sz
    elem(new double[s])  // allocate s doubles on the free
                        // store a pointer to those doubles in elem
{
    // Note: new does not initialize elements (but the standard vector
    // does)
}
```


- **new** allocates memory from the free store and returns a pointer to the allocated memory
- We use **new** to allocate objects that have to outlive the function that creates them



The free store

(sometimes called "the heap" or "dynamic memory")

- You request memory "to be allocated" "on the free store" by the **new** operator
 - The **new** operator returns a pointer to the allocated memory
 - A pointer is the address of the first byte of the memory
 - **int* p = new int;** *// allocate one uninitialized **int***
 *// **int*** means "pointer to **int**"*
 - **int* q = new int[7];** *// allocate seven uninitialized **ints***
 *// "an array of 7 **ints**"*
 - **double* pd = new double[n];** *// allocate **n** uninitialized **doubles***
 - A pointer points to an object of its specified type
 - A pointer does **not** know how many elements it points to



Access




- Individual elements

```
int* p1 = new int;           // get (allocate) a new
    uninitialized int
int* p2 = new int(5);        // get a new int initialized to 5

int x = *p2;                 // get/read the value pointed to by
    p2                       // (or "get the contents of what p2
    points to")
// in this case, the integer 5
int y = *p1;                 // undefined: y gets an undefined
    value: don't do that
```


Access



- Arrays (sequences of elements)

```
int* p3 = new int[5];           // get (allocate) 5 ints  
                                // array elements are numbered [0],  
                                [1], [2], ...
```

```
p3[0] = 7;                     // write to ("set") the 1st element  
    of p3
```

```
p3[1] = 9;
```


```
int x2 = p3[1];                 // get the value of the 2nd element of p3
```

```
int x3 = *p3;                  // we can also use the dereference  
    operator * for an array
```

```
// *p3 means p3[0] (and vice versa)
```

Pointer values

- Pointer values are memory addresses
 - Think of them as a kind of integer values
 - The first byte of memory is 0, the next 1, and so on




// you can see a pointer value (but you rarely need/want to):

```
int* p1 = new int(7);           // allocate an int and initialize it to 7
double* p2 = new double(7);      // allocate a double and initialize it to 7.0
cout << "p1==" << p1 << " *p1==" << *p1 << "\n"; // p1==??? *p1==c
cout << "p2==" << p2 << " *p2==" << *p2 << "\n"; // p2==??? *p2=7
```



Access

- A pointer does **not** know the number of elements that it's pointing to (only the address of the first element)

```
double* p1 = new double;  
*p1 = 7.3;           // ok  
p1[0] = 8.2;         // ok  
p1[17] = 9.4;         // ouch! Undetected error  
p1[-4] = 2.4;         // ouch! Another undetected error
```



```
double* p2 = new double[100];  
*p2 = 7.3;           // ok  
p2[17] = 9.4;         // ok  
p2[-4] = 2.4;         // ouch! Undetected error
```





- Fortunately, we have ways of avoiding such errors

Access


- A pointer does **not** know the number of elements that it's pointing to

```
double* p1 = new double;  
double* p2 = new double[100];
```



```
p1[17] = 9.4;    // error (obviously)
```

```
p1 = p2;         // assign the value of p2 to p1
```



(after the assignment)

```
p1[17] = 9.4;    // now ok: p1 now points to the array of  
100 doubles
```


Access

- A pointer **does** know the type of the object that it's pointing to

```
int* pi1 = new int(7);
```

```
int* pi2 = pi1; // ok: pi2 points to the same object as pi1
```


```
double* pd = pi1; // error: can't assign an int* to a double*
```

```
char* pc = pi1; // error: can't assign an int* to a char*
```

- There are no implicit conversions between a pointer to one value type to a pointer to another value type
- However, there are implicit conversions between *value* types:

```
*pc = 8; // ok: we can assign an int to a char
```

```
*pc = *pi1; // ok: we can assign an int to a char
```



The null pointer

- Sometimes, we need to say "this pointer doesn't point to anything just now"

```
double* p = nullptr;
```

```
// ...
```

```
If (p!=nullptr)           // p points to something
```

```
    *p = 7;
```

```
else
```

```
    *p = 9;           // No! never do this, p doesn't point to  
    anything
```

- More concisely, we can leave out the **!=nullptr**

```
If (p)           // p points to something, aka "p is valid"
```

```
    *p = 7;
```

- **nullptr** is commonly used to indicate

- End of a linked list

- No pointer value available just now

- No object to return a pointer to

A problem: memory leak

```
double* calc(int result_size, int max)
{
    double* p = new double[max];           // allocate max doubles on free
    store
    double* result = new double[result_size];
    // ... use p to calculate results to be put in result ...
    return result;
}

double* r = calc(200,100);    // oops!
```

- We “forgot” to give the memory allocated by **new** back to the free store
 - That doesn’t happen automatically (“no garbage collection”)
 - Lack of de-allocation (usually called “memory leaks”) can be a serious problem in real-world programs
 - A program that must run for a long time can’t afford any memory

We can give memory back to the free
store: **delete**

```
double* calc1(int result_size, int max)
{
    double* tmp = new double[max];    // allocate max doubles on free
    store

    double* result = new double[result_size];
    // ... use tmp to calculate results to be put in result ...

    delete[] tmp;                    // return the memory pointed to by
    tmp to the free store

    return result;
}

double* r = calc1(200,100);    // oops!

// ... use r ...

delete[ ] r;                    // return the memory pointed to by r to the
    free store
```


Memory leaks – resource leaks


- A program that needs to run “forever” can’t afford any memory leaks
 - An operating system is an example of a program that “runs forever”
- If a function leaks 8 bytes every time it is called, how many days can it run before it has leaked/lost a megabyte?
 - Trick question: not enough data to answer, but about 130,000 calls
- All memory is returned to the system at the end of the program
 - If you run using an operating system (Windows, Unix, whatever)
- Program that runs to completion with predictable memory usage may leak without causing problems
 - *i.e.*, memory leaks aren’t “good/bad” but they can be a major problem in specific circumstances
- Memory leaks is a major real-world problem
- Memory leaks is just a special case of resource leaks
 - E.g., files, locks, sockets

Memory leaks

- Another way to get a memory leak

```
void f()  
{  
    double* p = new double[27];  
    // ...  
    p = new double[42];  
    // ...  
    delete[] p;  
}
```

```
// 1st array (of 27 doubles)  
leaked
```



Memory leaks

- How do we systematically and simply avoid memory leaks?
 - Use resource handles
 - Use **vector**, etc.
 - don't mess directly with **new** and **delete**
 - Or use a garbage collector
 - A garbage collector is a program that keeps track of all of your allocations and returns unused free-store allocated memory to the free store
 - Not common in C++
 - not covered in this course; see <http://www.stroustrup.com/C++.html>
 - Unfortunately, a garbage collector doesn't prevent all resource leaks
 - Only memory leaks

A problem: memory leak

```
void f(int x)
{
    Vector v(x);           // define a Vector (which allocates x
    doubles on the free store)
    // ... use v ...

    // give the memory allocated by v back to the free store
    // but how? (Vector's elem data member is private)
}
```


Vector (destructor)

*// a very simplified **Vector** of **doubles**:*

```
class Vector {  
    int sz;                // the size  
    double* elem;          // a pointer to the elements  
public:  
    Vector(int s) :sz(s), elem(new double[s]) { }           // constructor:  
    allocates/acquires memory  
    ~Vector() { delete[ ] elem; }                            // destructor: de-  
    allocates/releases memory  
    // ...  
};
```

- Note: this is an example of a general and important technique:
 - acquire resources in a constructor
 - release them in the destructor
- Examples of resources: memory, files, locks, threads, sockets

Implicitly give memory back to the free store

```
Vector<double> calc2(int result_size, int max)
```

```
{
```

```
    Vector<double> tmp(max);
```

```
    // allocate max doubles
```

```
    on free store
```

```
    Vector<double> result(result_size);
```

```
    // ... use tmp to calculate results to be put in result ...
```

```
    return result;
```

```
} // tmp destroyed upon return
```

```
void user()
```

```
{
```

```
    // ...
```

```
    auto res = calc2(200,100); // oops!
```

```
    // ... use r ...
```

```
} // res destroyed upon return
```

Simpler and probably more efficient
than using **new** and **delete** explicitly
(yes, we can avoid copying a result
- Next lecture)

Free store summary

- Allocate using **new**

- New allocates an object on the free store, sometimes initializes it, and returns a pointer to it

- **int* pi = new int;** *// default initialization (none for int)*
 - **char* pc = new char('a');** *// explicit initialization*
 - **double* pd = new double[10];** *// allocation of (uninitialized) array*

- New throws a **bad_alloc** exception if it can't allocate (out of memory)

- Deallocate using **delete** and **delete[]**

- **delete** and **delete[]** return the memory of an object allocated by **new** to the free store so that the free store can use it for new allocations

- **delete pi;** *// deallocate an individual object*
 - **delete pc;** *// deallocate an individual object*
 - **delete[] pd;** *// deallocate an array*

- Delete of a zero-valued pointer ("the null pointer") does nothing

- **char* p = nullptr;**

Avoid “naked new” and “naked delete”

- Manual resource allocation and deallocation is error-prone
 - We forget to hand resources back (Ask any librarian)
 - We use pointers after the memory they point to has been deleted
 - Use of “raw pointers” to manage memory leads to overuse of pointers
- Using **vector** leads to simpler code
 - Compare `calc1()` and `calc2()`
 - See the following lectures and chapters

Generated destructors

- If a member of a class has a destructor, then that destructor will be called when the object containing the member is destroyed.

```
struct Customer {  
    string name;  
    vector<string> addresses;  
    // ...  
};  
  
void some_fct()  
{  
    Customer fred { "Fred", {"17 Oak St.", "232 Rock Ave."} };  
    // ... use fred ...  
}
```

- That saves us a lot of work
- And avoid a lot of bugs

Virtual destructors

- Destructors work correctly in class hierarchies
 - Provided you declare the destructor **virtual**

```
Shape* fct()
{
    Text tt {Point{200,200},"Anya"};           // local Text variable
    // ...
    return new Text{Point{100,100},"Courtney"}; // Text object
on the free store
}
```

```
void user()
{
    Shape* q = fct();
    // ... use the Shape without caring exactly which kind of shape it
    is ...
    delete q; // Shape's destructor is virtual so
Text::~~Text() is called if *q is a Text
```


But, what about “no naked **deletes**”?

- Use “resource-management” pointers ()

```
unique_ptr<Shape> fct()
{
    Text tt {Point{200,200},"Annemarie"};           // local Text variable
    // ...
    return make_unique<Text>(Point{100,100},"Nicholas"); // Text object on the free store
}

void user()
{
    unique_ptr<Shape> q = fct();
    // ... use the Shape without caring exactly which kind of shape it is ...
}
```

- Equivalent to the previous example, but simpler

Access to elements

- But our Vector doesn't have access operations

- So, let's add very simple ones

```
class Vector {                                // a very simplified vector of doubles  
public:  
    Vector(int s) :sz{s}, elem{new double[s]} { /* ... */ }           // constructor  
    ~Vector() { delete[] elem; }                                       // destructor  
  
    int size() const { return sz; }                                       // the current size  
  
    double get(int n) const { return elem[n]; }                       // access: read  
    void set(int n, double v) { elem[n]=v; }                         // access: write  
  
private:  
    int sz;                                // the size  
    double* elem;                          // a pointer to the elements  
};
```


Access to elements

- Not very elegant
 - But it'll do for now

```
Vector v(5);  
for (int i=0; i<v.size(); ++i) {  
    v.set(i,1.1*i);  
    cout << "v[" << i << "]==" << v.get(i) << '\n';  
}
```

Reminder

- Why look at the **Vector** implementation?
 - To see how the standard library **vector** really works
 - To introduce basic concepts and language features
 - Free store (heap)
 - Copying
 - Dynamically growing data structures
 - To see how to directly deal with memory
 - To introduce the techniques and concepts we need to understand C
 - Including the dangerous ones (and see how to avoid those in C++)
 - To demonstrate class design techniques
 - To see examples of “neat” code and good design

Next lecture

- We'll see how we can change our **Vector**'s implementation to better allow for changes in the number of elements. Then we'll modify **Vector** to take elements of an arbitrary type and add range checking. That'll imply looking at templates and revisiting exceptions.