


Chapter 19 – containers and iterators

*Any problem in computer science can be solved
with another layer of indirection.
Except, of course,
the problem of too many indirections.
– David J. Wheeler*



Abstract

- This lecture and the two following present the STL
 - the containers and algorithms part of the C++ standard library
- The STL is an extensible framework dealing with data in a C++ program.
 - the general ideal
 - the fundamental concepts
 - examples of containers and algorithms.
- Key notions to tie data together with algorithms (for general processing)
 - *sequence* (aka *range*)
 - iterator

Common tasks

- Collect data into containers
- Organize data
 - For printing
 - For fast access
- Retrieve data items
 - By index (e.g., get the Nth element)
 - By value (e.g., get the first element with the value "Chocolate")
 - By properties (e.g., get the first elements where **age**<64)
- Add data
- Remove data
- Sorting and searching
- Simple numeric operations

Lifting example (concrete algorithms)

- Jack and Jill each deliver some data
 - Jack in traditional (old-fashioned) C style
 - Jill in contemporary C++ style

```
double* get_from_jack(int* count);
```

*// Jack fills an array and puts the number of elements in *count*

```
vector<double> get_from_jill();
```

// Jill fills a vector

```
void fct()
```

```
{
```

```
    int jack_count = 0;
```

```
    double* jack_data = get_from_jack(&jack_count);
```

```
    vector<double> jill_data = get_from_jill();
```

```
    // ... process ...
```

```
    delete[] jack_data;
```

```
}
```


Simple use ("processing")

```
double h = -1;
double* jack_high;           // jack_high will point to the element with the highest value
double* jill_high;           // jill_high will point to the element with the highest value

for (int i=0; i<jack_count; ++i)
    if (h<jack_data[i]) {
        jack_high = &jack_data[i];           // save address of largest element
        h = jack_data[i];                     // update "largest element"
    }

h = -1;
for (double& x : jill_data)
    if (h<x) {
        jill_high = &x;                       // save address of largest element
        h = x;                                 // update "largest element"
    }

cout << "Jill's max: " << *jill_high
      << "; Jack's max: " << *jack_high;
```


Similar,
but different in most details

A first attempt to unify/generalize

```
double* high(double* first, double* last)
```

// return a pointer to the element in [first:last) that has the highest value

```
{  
    double h = -1;  
    double* high;  
    for (double* p = first; p!=last; ++p)  
        if (h<*p) {  
            high = p;  
            h = *p;  
        }  
    return high;  
}
```




```
double* jack_high = high(jack_data, jack_data+jack_count);  
double* jill_high = high(&jill_data[0], &jill_data[0]+jill_data.size());
```

```
cout << "Jill's max: " << *jill_high  
      << "; Jack's max: " << *jack_high;
```

Beware:
We left two errors behind

We “accidentally” generalized

- We can now do `high()` for parts of an array



```
vector<double>& v = *jill_data;  
double* middle = &v[0]+v.size()/2;  
double* high1 = high(&v[0], middle);           // max of first half  
double* high2 = high(middle, &v[0]+v.size());  // max of second half
```

Ideals

We'd like to write common programming tasks so that we don't have to re-do all the work each time we find

- a new way of storing the data
- A slightly different way of representing the data
- a slightly different way of interpreting the data
- Some different processing to do
- Observations
 - Using an **int** isn't that different from using a **double**
 - Using a **vector<int>** isn't that different from using a **vector<string>**
 - Finding a value in a **vector** isn't all that different from finding a value in a **list** or an array
 - Looking for a **string** ignoring case isn't all that different from looking at a **string** not ignoring case
 - Graphing experimental data with exact values isn't all that different from graphing rounded values
 - Copying a file isn't all that different from copying a vector

Ideals (continued)

- Code that's
 - Easy to read
 - Easy to modify
 - Regular
 - Short
 - Fast
- Uniform access to data
 - Independently of how it is stored
 - Independently of its type
- ...

Ideals (continued)

- ...
- Type-safe access to data
- Easy traversal of data
- Compact storage of data
- Fast
 - Retrieval of data
 - Addition of data
 - Deletion of data
- Standard versions of the most common algorithms
 - Copy, find, search, sort, sum, ...

Examples

- Sort a vector of strings
- Find a number in a phone book, given a name
- Find the highest temperature
- Find all values larger than 800
- Find the first occurrence of the value 17
- Sort the telemetry records by unit number
- Sort the telemetry records by time stamp
- Find the first value larger than “Petersen”?
- What is the largest amount seen?
- Find the first difference between two sequences
- Compute the pairwise product of the elements of two sequences
- What are the highest temperatures for each day in a month?
- What are the top 10 best-sellers?
- What’s the entry for “C++” (say, in Google)?
- What’s the sum of the elements?

Generic programming


- Start with a concrete algorithm
 - Or better yet: a set of related concrete uses
- Generalize it until it makes the minimal assumptions needed
 - Without losing performance
- That's sometimes called “lifting an algorithm”
 - We go from the concrete to the more abstract
 - The other way most often leads to bloat
 - We are concerned with performance
 - Slow code will eventually be thrown away
 - Our aim (for the end user) is
 - Greater range of uses (re-use)
 - More correctness
 - Through better specification

The STL

- Part of the ISO C++ Standard Library
 - About 100 algorithms
 - About 12 containers
- Mostly non-numerical
 - Only a few standard algorithms specifically do numerical computation
 - Accumulate(), inner_product(), partial_sum(), adjacent_difference()
 - Handles textual data as well as numeric data
 - E.g. string
 - Deals with organization of code and data
 - Built-in types, user-defined types, and data structures
- Optimizing disk access was among its original uses
 - Performance was always a key concern

Basic model

- A pair of iterators defines a sequence
 - The beginning (points to the first element - if any)
 - The end (points to the one-beyond-the-last element)




- An iterator is a type that supports the “iterator operations”
 - ++ Go to next element
 - * Get value
 - == Does this iterator point to the same element as that iterator?
- Some iterators support more operations (e.g. --, +, and [])

Basic model


- Algorithms

sort, find, search, copy, ...



- Containers

vector, list, map, unordered_map, ...



- Separation of concerns

- Algorithms manipulate data, but don't know about containers
- Containers store data, but don't know about algorithms
- Algorithms and containers interact through iterators
 - Each container has its own iterator types

Back to Jack and Jill

- Converting “Jack and Jill” to STL style

```
template<forward_iterator Iter>
```

```
Iter high(Iter first, Iter last)
```

```
    // return an iterator to the element in [first:last) that has the highest value
```

```
{
```

```
    Iter high = first;
```

```
    for (Iter p = first; p!=last; ++p)
```

```
        if (*high<*p)
```

```
            high = p;
```

```
    return high;
```

```
}
```

We still left an error behind.
Now is a good time to find it.


How did the other one disappear?

Back to Jack and Jill


```
double* get_from_jack(int* count);           // Jack fills an array and puts the number of elements in  
    *count  
vector<double> get_from_jill();              // Jill fills the vector  
  
void fct()  
{  
    int jack_count = 0;  
    double* jack_data = get_from_jack(&jack_count);  
    vector<double> jill_data = get_from_jill();  
  
    double* jack_high = high(jack_data, jack_data + jack_count);  
    double* jill_high = high(jill_data.begin(), jill_data.end());  
  
    cout << "Jill's high " << *jill_high << "; Jack's high " << *jack_high;  
    delete[] jack_data;  
}
```

STL-style vector and list


- The general model of a sequence



- **vector**



- **list** (a doubly linked)



An STL-style list

```
template<Element T>
class List {
    // ... representation and implementation details ...
public:
    // ... constructors, destructor, etc. ...

    class iterator;                // member type: iterator
    iterator begin();              // iterator to first element
    iterator end( );              // iterator to one beyond last
    element

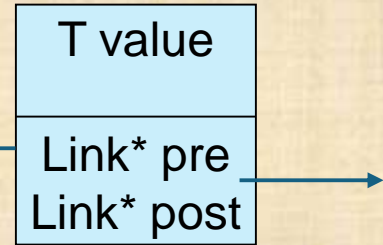
    iterator insert(iterator p, const T& v);    // insert v into list
    after p

    iterator erase(iterator p);                // remove p from the list

    void push_back(const T& v);                // insert v at end
    void push_front(const T& v);              // insert v at front
    void pop_front();                          // remove the first element
    void pop_back();                          // remove the last element

    T& front();                               // the first element
    T& back();                               // the last element
};
```

Link:



Iteration - a List iterator

```
template<Element T>
class List<Elem>::iterator {
    Link<T>* curr;                // current link

public:
    iterator(Link<T>* p) :curr{p} { }

    iterator& operator++() {curr = curr->succ; return *this; }
    // forward

    iterator& operator--() { curr = curr->prev; return *this; }
    // backward

    T& operator*() { return curr->val; }                // get
value (dereference)

    bool operator==(const iterator& b) const { return curr==b.curr; }
    bool operator!=(const iterator& b) const { return curr!=b.curr; }

};
```


“Jack and Jill” with a list


```
void f()
{
    list<int> lst;
    for (int x; cin >> x; )           // build a list from
input
    lst.push_front(x);

    list<int>::iterator p = high(lst.begin(), lst.end()); // traverse the
list to find the highest element
    cout << "the highest value was " << *p << '\n';
}
```


But what if the input was empty

Traversal ("empty" isn't a special case)

General model of a sequence:



Empty sequence:



- "Jack and Jill" again

```
list<int>::iterator p = high(lst.begin(), lst.end());  
if (p==lst.end())           // did we reach the end?  
    cout << "The list is empty";  
else  
    cout << "the highest value is " << *p << '\n';
```


Vector member types

```
template<Element T, Allocator A = allocator<T>>
class Vector {
public:
    using size_type = int;           // number of elements
    using value_type = T;           // type of an element
    using iterator = T*;            // type of an iterator
    using const_iterator = const T*; // type of an iterator to
    elements you can't modify
    // ...
    iterator begin();
    const_iterator begin() const;
    iterator end();
    const_iterator end() const;

    size_type size();
    // ...
```

Vector traversal

- Old-fashioned use of size and subscript
 - We can get the size wrong
 - We can get the type of the loop variable wrong
 - Range-checking can be costly

```
void print1(const vector<double>& v)
{
    for (int i = 0; i<v.size(); ++i)
        cout << v[i] << '\n';
}
```


Container traversal (works for all STL containers)

- Using iterators
 - Verbose
 - We can (still) get the range wrong, though less likely that with indexes

```
void print2(const vector<double>& v, const list<double>& lst)
{
    for (vector<T>::iterator p = v.begin(); p!=v.end(); ++p)
        cout << *p << '\n';

    for (list<T>::iterator p = lst.begin(); p!=lst.end(); ++p)
        cout << *p << '\n';
}
```

Container traversal (works for all STL containers)

- Use **auto** to get the iterator type
 - We can (still) get the range wrong, though less likely that with indexes

```
void print3(const vector<double>& v, const list<double>& lst)
{
    for (auto p = v.begin(); p!=v.end(); ++p)
        cout << *p << '\n';

    for (auto p = lst.begin(); p!=lst.end(); ++p)
        cout << *p << '\n';
}
```


Container traversal (works for all STL containers)


- Use a **range**-for loop
 - We can't get the range wrong
 - Range checking is easy
 - Prefer a **range**-for (or a range algorithm; see chapter 21)

```
void print4(const vector<double>& v, const list<double>& lst)
{
    for (double x : v)
        cout << x << '\n';


    for (double x : lst)
        cout << x << '\n';
}
```

`insert()` into vector

```
vector<int>::iterator p =  
    v.begin();  
++p; ++p; ++p;  
vector<int>::iterator q = p;  
++q;
```




```
p=v.insert(p,99);    // leaves p pointing at the inserted  
    element
```




- Note: `q` is invalid after the `insert()`
- Note: Some elements moved; all elements could have moved

erase() from vector




```
p = v.erase(p);    // leaves p pointing at the element after  
the erased one
```




- vector elements move when you `insert()` or `erase()`
- Iterators into a vector are invalidated by `insert()` and `erase()`

`insert()` into list

```
list<int>::iterator p = v.begin(); ++p; ++p; ++p;  
list<int>::iterator q = p; ++q;
```




```
v = v.insert(p, 99);    // leaves p pointing at the inserted  
element
```




- Note: `q` is unaffected
- Note: No elements moved around

`erase()` from list



`p = v.erase(p);` // leaves `p` pointing at the element after the
erased one



- Note: list elements do not move when you `insert()` or `erase()`

Implementing **Vector** **erase()** and **insert()**

- Still working to complete our Vector
 - Prefer **std::vector**
 - Remember the **Vector** layout



Vector::erase()

- With a little help from the standard library implementations are getting manageable

```
template<Element T, Allocator A>
Vector<T,A>::iterator Vector<T,A>::erase(iterator p)
{
    if (p==end())
        return p;
    move(p+1,r.sz,p);          // move each element one position to the
left
    destroy_at(r.elem()+r.sz-1)); // destroy surplus last element
    --r.sz;
    return p;
}
```

Vector::insert()

- With a little help from the standard library implementations are getting manageable

```
template<Element T, Allocator A>
Vector<T,A>::iterator Vector<T,A>::insert(iterator p, const T& val)
{
    int index = p-begin(); // save index in case of
    relocation
    if (size()==capacity())
        reserve(size()==0?8:2*size()); // make sure we have space
    p = begin()+i; // p now points into the current
    allocation
    move_backward(p,r.sz-1,p+1); // move each element one position
    to the right
    *(begin()+index) = val; // `insert' val
    ++r.sz;
```


vector, list, and string

- By default, use a **vector**
 - You need a reason not to
 - You can “grow” a vector (e.g., using **push_back()**)
 - You can **insert()** and **erase()** in a vector
 - Vector elements are compactly stored and contiguous
 - For small vectors of small elements all operations are fast
 - compared to lists
- If you don’t want elements to move, use a **list**
 - You can “grow” a list (e.g., using **push_back()** and **push_front()**)
 - You can **insert()** and **erase()** in a list
 - List elements are separately allocated

vector, list, and string

- Use a **string** when you are doing string operations
 - Elements are characters (you can select the character set)
 - **strings** have concatenation (+ and +=)
 - **strings** are expandable and mutable
 - Small **strings** do not use free store
- Note that there are more containers (see chapter 20), e.g.,
 - **map**
 - **unordered_map**
 - **set**
 - **unordered_set**

Next lecture

- Map (aka dictionaries or associative arrays), unordered_map (aka tables), and sets