


Chapter 20 – Maps and Sets



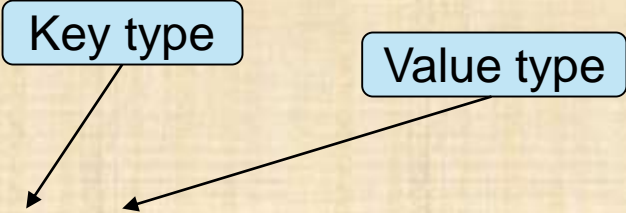
*Write programs that do one
thing
and do it well.
Write programs to work
together.
– Doug McIlroy*

Overview

- Maps and unordered_maps
 - Balanced trees and hash tables
- Timing
- Sets
- “Almost containers”
 - Adapting container-like data structures to the STL

Map (an associative array)

- For a **vector**, you subscript using an integer
- For a **map**, you can define the subscript to be any type with an order



```
int main()
{
    map<string,int> words;           // keep (word,frequency) pairs
    for (string s; cin>>s; )        // words is subscripted by a string
        ++words[s];                // string < (less than) determines the order
                                    // words[s] returns an int&
                                    // the int values are initialized to 0

    for (const auto& [key,value] : words)
        cout << key << ": " << value << "\n";
}
```

The words program (word frequencies)

- Input and output

C++ is a general purpose programming language designed to make programming more enjoyable for the serious programmer.

Except for minor details, C++ is a superset of the C programming language.

In addition to the facilities provided by C, C++ provides flexible and efficient facilities for defining new types.

C: 1

C++: 3

C,: 1

Except: 1

In: 1

a: 2

addition: 1

and: 1

by: 1

defining: 1

designed: 1

details,: 1

efficient: 1

enjoyable: 1

facilities: 2

flexible: 1

for: 3

general: 1

more: 1

is: 2

language: 1

language.: 1

make: 1

minor: 1

new: 1

of: 1

programmer.: 1

programming: 3

provided: 1

provides: 1

purpose: 1

serious: 1

superset: 1

the: 3

to: 2

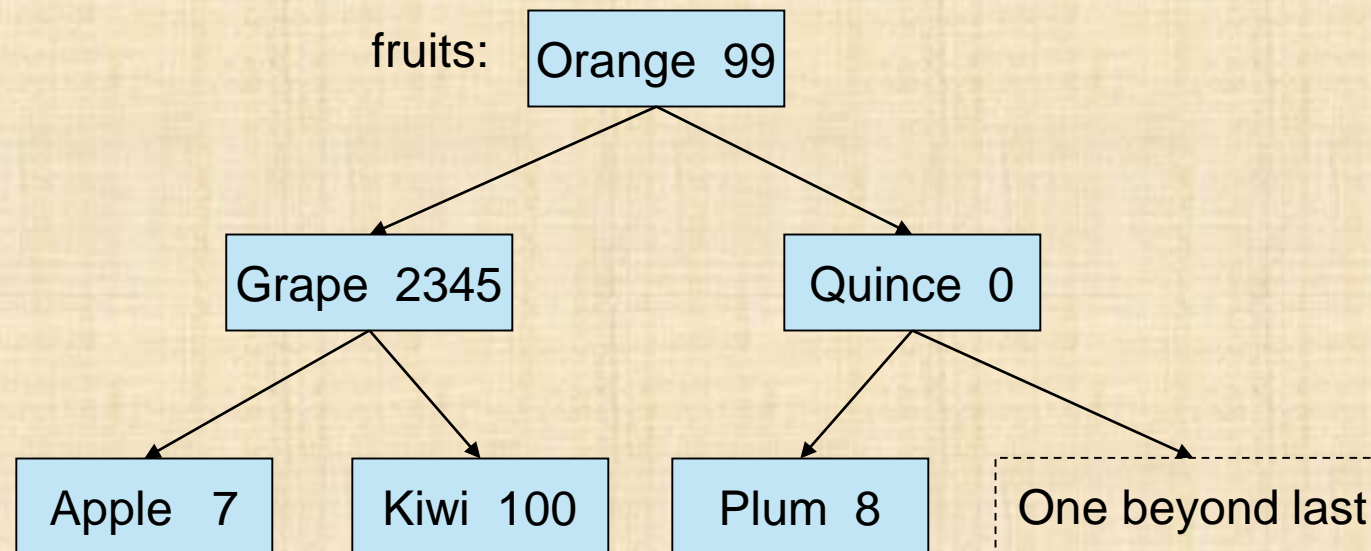
types.: 1

Map

Map node:

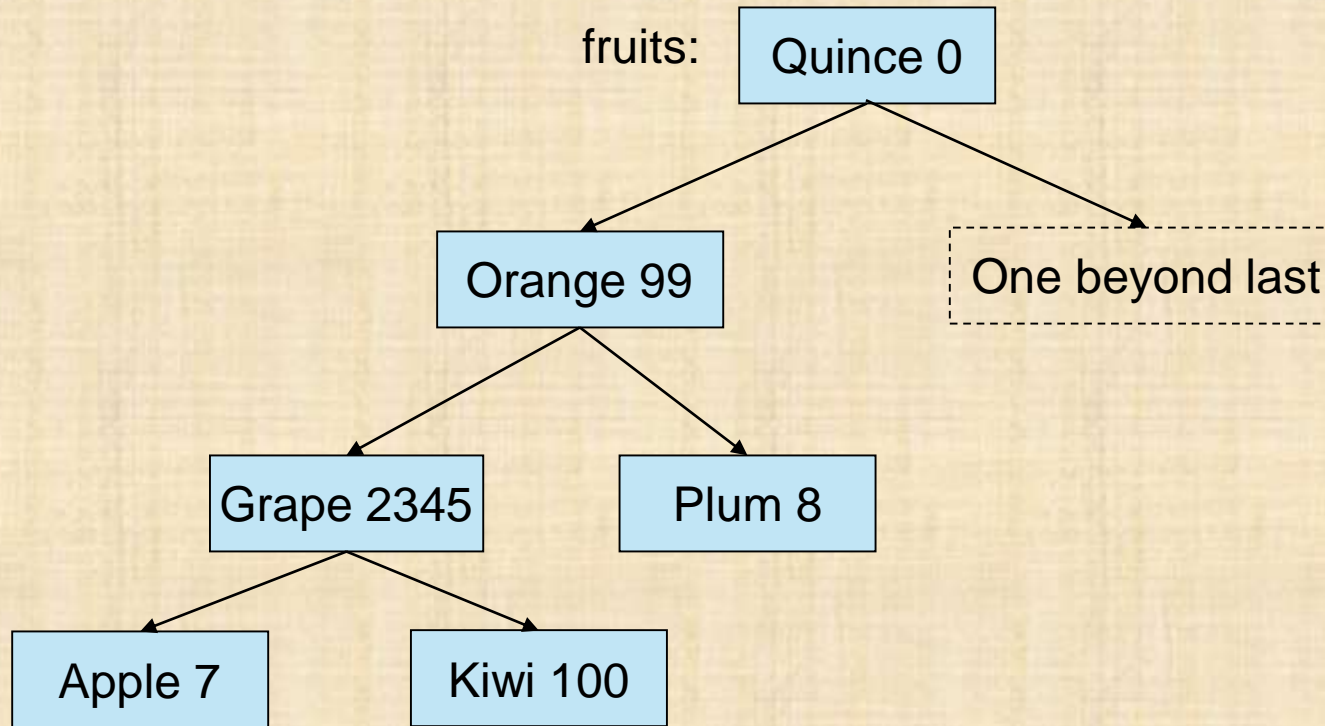
Key first Value second
Node* left Node* right ...

- After **vector**, **map** is the most useful standard library container
 - Maps (and/or hash tables) are the backbone of scripting languages
- A **map** is really an ordered balanced binary tree
 - By default, ordered by < (less than)
 - `map<Fruit,int> fruits = { {Kiwi,100}, {Quince,0}, {Plum,8}, {Apple,7}, {Grape,2345}, {Orange,99} };`



Map

- Maps are balanced
 - Giving on average $\log_2(N)$ indirections to reach a node
- Unbalanced trees often require more indirections
 - `map<Fruit,int> fruits = { {Kiwi,100}, {Quince,0}, {Plum,8}, {Apple,7}, {Grape,2345}, {Orange,99} };`



Some implementation defined type

Map

```
template<class Key, class Value>
class map { // note the similarity to vector and list
    // ...
    using value_type = pair<Key, Value>; // a map deals in (Key, Value) pairs

    using iterator = ???; // probably a pointer to a tree
    node

    using const_iterator = ???;

    iterator begin(); // points to first element
    iterator end(); // points to one beyond the last element

    Value& operator[ ](const Key&); // get Value for Key; creates pair if
    // necessary, using Value( )

    iterator find(const Key& k); // is there an entry for k?

    void erase(iterator p); // remove element pointed to by p

    pair<iterator, bool> insert(const value_type&); // insert new (Key, Value)
    pair; the bool is false if insert failed
```

Map example (build some maps)

```
// Values from www.djindexes.com
```

```
map<string,double> dow_price = {           // Dow Jones Industrial index  
    (symbol,price)  
    {"MMM",104.48}, {"AAPL",165.02}, {"MSFT",285.76},  
    // ...  
};  
  
map<string,double> dow_weight = { // Dow (symbol,weight)  
    {"MMM", 2.41}, {"AAPL",2.84}, {"MSFT",4.88},  
    // ...  
};  
  
map<string,string> dow_name = {           // Dow (symbol,name)  
    {"MMM","3M"}, {"AAPL","Apple"}, {"MSFT","Microsoft"},  
    // ...  
};
```


Map example (some uses)

```
double caterpillar = dow_price ["CAT"];           // read values from
a map
double boeing_price = dow_price ["BA"];

if (dow_price.find("INTC") != dow_price.end())    // find an entry in
a map
    cout << "Intel is in the Dow\n";

// Iterating through a map is easy:
for (const auto& [symbol,price] : dow_price)      // output names in
alphabetic order of symbols
    cout << symbol << '\t' << price << '\t' << dow_name[symbol] << '\n';
```

Unordered_map (aka hash table)

- An unordered_map is like a map, except
 - Lookup is by a "hash function"
 - An **unordered_map**'s elements are not in order

```
unordered_map<string,double> dow_price= {           // Dow Jones Industrial
    index (symbol,price)
    {"MMM",104.48}, {"AAPL",165.02}, {"MSFT",285.76},
    // ...
};
```

```
for (const auto& [symbol,price] : dow_price)           // output names in no
defined order
    cout << symbol << '\t' << price << '\t' << dow_name[symbol] <<
'\n';
```


Lookup: vector, map, and unordered_map


- Linear search in a vector is proportional to N , the number of elements
 - But each lookup is very cheap
- Map Lookup cost is proportional to $\log_2(N)$
 - Indirections through pointers
- Unordered_map lookup cost is proportional to 1 (constant time)

vector	N	16	128	1024	$16 \cdot 1024$	$1024 \cdot 1024$	$1024 \cdot 1024 \cdot 1024$	ue
map	$\log_2(N)$	4	7	10	14	20	30	
unordered_map	1	1	1	1	1	1	1	


- The performance of an unordered_map is critically dependent on the hash function

Lookup: vector, map, and unordered map


- Vector (linear search)



- Map (tree walk):



- Unordered_map (hash lookup):



Timing

- Never talk about “efficiency” without measurement
 - Space
 - Runtime
 - Compile time
- Beware of
 - Tiny examples
 - Useful complexities that only make sense for huge examples
- Don't tryst your “intuition”/guesses
 - Even experts can guess wrong by a factor of 1000
- Always run a test at least three times
 - And test for consistency
- Be suspicious of your results
 - Always try to explain them

Timing

- Simple tests are much better than mere guesses

```
using namespace chrono;           // that's where the timing support is

auto t0 = system_clock::now();    // the point of time of the call
auto x = do_something();
auto t1 = system_clock::now();
cout << "res: " << x << '\n';

cout << t1-t0 << '\n';           // that's how long it took (default
    unit)

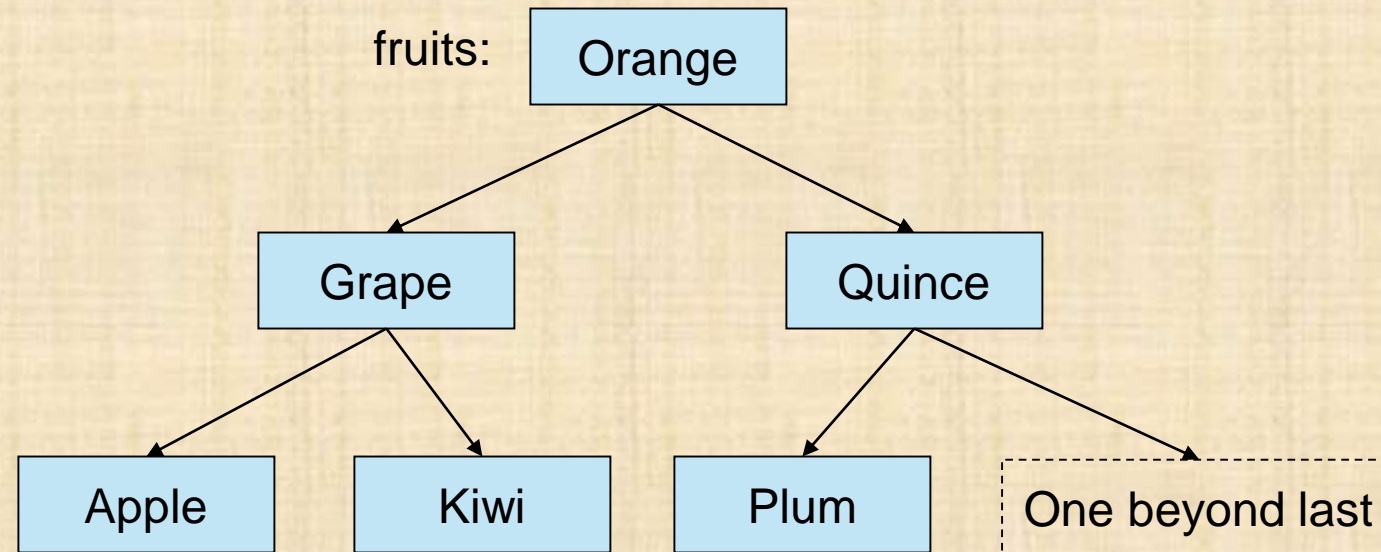
cout << duration_cast<microseconds>(t1-t0).count() << "us\n"; // that's
    how long it took in useconds
```


Set

set node:

Key first
Node* left Node* right ...

- A **set** is really an ordered balanced binary tree
 - By default, ordered by <
 - `set<string> fruits = { Kiwi, Quince, Plum, Apple, Grape, Orange };`



Set

- We can use a set to order (sort) values

```
int main()
{
    string from, to;
    cin >> from >> to;           // get source and target file names

    ifstream is {from};           // make input stream
    ofstream os {to};             // make output stream

    set<string> b {istream_iterator<string>{is},
istream_iterator<string>{}};      // read into the set
    copy(b.begin() ,b.end() , ostream_iterator<string>{os, "\n"});
    // copy to output
}
```


The sorting program (using a set)

- Input and output

C++ is a general purpose programming language designed to make programming more enjoyable for the serious programmer.

Except for minor details, C++ is a superset of the C programming language.

In addition to the facilities provided by C, C++ provides flexible and efficient facilities for defining new types.

C++
C,
Except
In
a
addition
and
by
defining
designed
details,
efficient
enjoyable
facilities
flexible
for
general
more
is
language
language.
make
minor
new
of
programmer.
programming
provided
provides
purpose
serious
superset
the
to
types.

A container (incomplete overview)

- An STL container is characterized by having a set of useful types and operations
- Has a sequence of elements (**begin()**:**end()**)
- Provides copy operations that copy all elements.
- Provides move operations that move all elements
- Names its element type **value_type**
- Has iterator types called **iterator** and **const_iterator**
- Iterators provide *****, **++** (both prefix and postfix), **==**, and **=** with the appropriate semantics
- The iterators for **list** also provide **--** for moving backward; that's called a bidirectional iterator
- The iterators for vector also provide **--**, **[]**, **+**, and **-** and are called random-access iterators
- Provides **insert()** and **erase()**, **front()**, and **back()**, **push_back()**, and **pop_back()**, **size()**, **swap()**, etc.
- **vector** and **map** also provide subscripting (e.g., operator **[]**)
- Provides comparison operators (**==**, **!=**, **<**, **<=**, **>**, and **>=**) that compare the

Standard containers

- **vector** a contiguously allocated sequence of elements; use it as the default container
- **list** a doubly-linked list; use to insert and delete elements without moving existing elements
- **forward_list** a singly-linked list; use for lists that are mostly empty
- **deque** a cross between a **list** and a **vector**
don't use until you have expert-level knowledge of algorithms and machine architecture
- **map** a balanced ordered tree; use it when you need to access elements by value (§ 20.2)
- **multimap** a balanced ordered tree where there can be multiple copies of a key
- **set** a balanced ordered tree; use it when you need to keep track of individual values (§ 20.5)
- **multiset** a balanced ordered tree where there can be multiple copies of a key
- **unordered_map** a hash table; an optimized version of **map**;
use for large maps when you need high performance
and can devise a good hash function (§20.3)

“Almost containers”

- Data types that provide much of what is required from a standard container, but not all
 - There are many because needs of data representations are many
 - Most are not in the standard (but those may be essential for an application)
- **T[n]** a built-in array; no **size()** or other member functions; prefer a container, such as **vector**, **string**, or **array**, over a built-in array when you have a choice
- **array** a fixed-size array that doesn't suffer most of the problems related to the built-in arrays
- **string** holds only characters but provides operations useful for text manipulation, such as concatenation (**+** and **+=**); prefer the standard string to other strings
- **valarray** a numerical vector with mathematical vector operations, but with many restrictions to encourage high-performance implementations; use only if you do a lot of vector arithmetic

Built-in array and **array**

- You can use **std::array** much like a container
 - A built-in array doesn't carry its number of elements out of its scope; a **std::array** does
 - A built-in implicitly converts to a pointer ("forgetting its size"); a **std::array** does not
 - Identical size and speed

```
void array_test(auto& a1, auto a2)    // or f(array<double,6>& , double
arr2[] )
{
    ranges::sort(a1);                 // OK
    sort(begin(a1), end(a1));         // OK

    ranges::sort(a2);                 // error: we don't know where the end
of the array is
    sort(begin(a2), end(a2));         // error: we don't know where the end
of the array is
}
```

```
array<double, 6> arr1 = { 0, 0, 1, 1, 3, 3, 5, 5, 2, 2, 4, 4 };
```

Input and output iterators

- we can provide iterators for **istreams** and **ostreams**
 - Thus, make them usable like containers

```
ostream_iterator<string> oo(cout); // assigning to *oo is to
write to cout
*oo = "Hello, ";                // meaning cout << "Hello, "
++oo;                          // "get ready for next output
operation"
*oo = "world!\n";              // meaning cout << "world!\n"

istream_iterator<string> ii(cin); // reading *ii is to read a
string from cin
string s1 = *ii;               // meaning cin>>s1
++ii;                          // "get ready for the next
input operation"
string s2 = *ii;               // meaning cin>>s2
```


Make a quick dictionary (using a vector)

- Now we can treat **istreams** and **ostreams** as if they were containers
 - As we say: "copy from input" and "copy to output"

```
istream_iterator<string> ii(cin);           // make input iterator
for cin

istream_iterator<string> eos;               // input sentinel (defaults
to EOF)

ostream_iterator<string> oo(cout, "\n");    // make output iterator for
cout; append "\n" each time

vector<string> buf {ii, eos};                // buf is a vector initialized
from input

ranges::sort(buf);                          // sort the buffer

ranges::unique_copy(buf, oo);               // copy buffer cout; discard
replicated values
```

The sorting program (using a vector_c)

- Input and output

C++ is a general purpose programming language designed to make programming more enjoyable for the serious programmer.

Except for minor details, C++ is a superset of the C programming language.

In addition to the facilities provided by C, C++ provides flexible and efficient facilities for defining new types.

C++
C,
Except
In
a
addition
and
by
defining
designed
details,
efficient
enjoyable
facilities
flexible
for
general

more
is
language
language.
make
minor
new
of
programmer.
programming
provided
provides
purpose
serious
superset
the
to
types.

Which version do you think is faster?

The one using a map or the one using a vector?

Why?

Range checking

- But what if we were writing into a fixed-sized buffer?

- Buffer overflows are nasty

- Consider

```
vector<int> v = {0,1,2,3,4,5};
```

```
vector<int> v1(10);
```

```
vector<int> v2(5);
```

- #1 check the size

```
If (v.size() < v1.size()) Ranges::copy(v,v1); else throw  
Range_error{};
```

```
If (v.size() < v2.size()) Ranges::copy(v,v2); else throw  
Range_error{};
```

- #2 use a checked iterator

```
ranges::copy(v,Output_range{v1});
```

```
// copies v into v1
```

```
ranges::copy(v,Output_range{v2});
```

```
// throws
```

```
Range_error
```

Range checking

- But what if we were writing into a fixed-sized buffer?

```
template<ranges::range R>
class Output_range {
public:
    using value_type = ranges::range_value_t<R>;
    using difference_type = int;
    Output_range(R r) : b{ r.begin() }, e{ r.end() }, p{ b } {}
    Output_range& operator++() { check_end(); ++p; return *this;
}

    Output_range operator++(int) { check_end(); auto t{ *this };
    ++p; return t; }
    value_type& operator*() const { check_end(); return *p; }
private:
    void check_end() const { if (p == e) throw Range_error{}; }
    ranges::iterator_t<R> b;
    ranges::iterator_t<R> e;
    ranges::iterator_t<R> p;
```


Ranges and iterators

- Iterators can be a source of errors

```
sort(v.end(), v.start());           // Oops!  
sort(v1.start(), v2.end());         // Oops!
```

- Prefer ranges

- They more precisely say what's intended

```
ranges::sort(v);                     // if this doesn't make sense the  
compiler will tell us
```

- we can define a range in three ways:

- *{begin, end}* a pair of iterators
 - *{begin, length}* an iterator and a number of elements
 - *{begin, predicate}* an iterator and predicate to determine if the end has been reached

Ranges and iterators

- Iterator categories

- **input iterator** Can iterate forward using **++** and read element values using *****.
This is the kind of iterator that **istream** offers.
If **(*p).m** is valid, **p->m** can be used as a shorthand.
- **output iterator** Can iterate forward using **++** and write element values using *****.
This is the kind of iterator that **ostream** offers.
- **forward iterator** An input iterator that can iterate repeatedly over a sequence
and repeatedly read from or write to an element.
This is the kind of iterator that **forward_list** offers.
- **bidirectional iterator** A forward iterator that can move backward (using **--**).
This is the kind of iterator that **list**, **map**, and **set** offer.
- **random-access iterator** A bidirectional iterator that can move forward

Next lecture

- Algorithms!